

Natural Language Processing for document classification

Abstract

In this study we will apply support vector machines to the domain of natural language processing (NLP) for the purposes of sentiment analysis. Our approach will be to use support vector machines to automatically classify many text documents into mutually exclusive groups. Note that this is a *supervised* learning technique.

1. Supervised Document Classification

Consider a collection of text documents. Each document has an associated set of keywords, which we will call features. Each of these documents might possess a class label describing what the article is about.

For instance, a set of articles from a website discussing pets might have articles that are primarily about dogs, cats or hamsters (say). Certain words, such as "cage" (hamster), "leash" (dog) or "milk" (cat) might be more representative of certain pets than others. Supervised classifiers are able to isolate certain words which are representative of certain labels by learning from a set of training articles that are already pre-labelled, often in a manual fashion by a human.

Mathematically each of the j articles about pets within a training corpus have an associated feature vector x_j , with components of this vector representing strength of words (we will define strength below). Each article also has an associated class label, y_j , which in this case would be the name of the pet most associated with the article. The classifier is fed feature vector-class label pairs and it learns how representative features are for particular class labels. In the following new example we will use the SVM as our model and train it on a corpus (a collection of documents) that has been previously generated.

You will see that all the files beginning with **reut2-** are **.sgm**, which means that they are Standard Generalized Markup Language(SGML)[1] files. Unfortunately, Python deprecated **sgmlib** from Python in 2.6 and fully removed it for Python 3. However, all is not lost because we can create our own SGML Parser class that overrides Python's built in **HTMLParser**[2].

Here is a single news item from one of the files:

```
..
..
<REUTERS TOPICS="YES" LEWISSPLIT="TRAIN"
CGISPLIT="TRAINING-SET" OLDDID="5544" NEWID="1">
<DATE>26-FEB-1987 15:01:01.79</DATE>
<TOPICS><D>cocoa</D></TOPICS>
```

<PLACES><D>el-salvador</D><D>usa</D><D>uruguay</D></PLACES>

<PEOPLE></PEOPLE>

<ORGS></ORGS>

<EXCHANGES></EXCHANGES>

<COMPANIES></COMPANIES>

<UNKNOWN>

C T

f0704reute

u f BC-BAHIA-COCOA-REVIEW 02-26 0105</UNKNOWN>

<TEXT>

<TITLE>BAHIA COCOA REVIEW</TITLE>

<DATELINE> SALVADOR, Feb 26 - </DATELINE><BODY>

Showers continued throughout the week in the Bahia cocoa zone, alleviating the drought since early January and improving prospects for the coming temporao, although normal humidity levels have not been restored, Comissaria Smith said in its weekly review. The dry period means the temporao will be late this year. Arrivals for the week ended February 22 were 155,221 bags of 60 kilos making a cumulative total for the season of 5.93 mln against 5.81 at the same stage last year. Again it seems that cocoa delivered earlier on consignment was included in the arrivals figures.

Comissaria Smith said there is still some doubt as to how much old crop cocoa is still available as harvesting has practically come to an end. With total Bahia crop estimates around 6.4 mln bags and sales standing at almost 6.2 mln there are a few hundred thousand bags still in the hands of farmers, middlemen, exporters and processors.

There are doubts as to how much of this cocoa would be fit for export as shippers are now experiencing difficulties in obtaining +Bahia superior+ certificates.

In view of the lower quality over recent weeks farmers have sold a good part of their cocoa held on consignment.

Comissaria Smith said spot bean prices rose to 340 to 350 cruzados per arroba of 15 kilos.

Bean shippers were reluctant to offer nearby shipment and only limited sales were booked for March shipment at 1,750 to 1,780 dlrs per tonne to ports to be named.

New crop sales were also light and all to open ports with June/July going at 1,850 and 1,880 dlrs and at 35 and 45 dlrs under New York july, Aug/Sept at 1,870, 1,875 and 1,880 dlrs per tonne FOB.

Routine sales of butter were made. March/April sold at

4,340, 4,345 and 4,350 dlrs.
April/May butter went at 2.27 times New York May, June/July at 4,400 and 4,415 dlrs, Aug/Sept at 4,351 to 4,450 dlrs and at 2.27 and 2.28 times New York Sept and Oct/Dec at 4,480 dlrs and 2.27 times New York Dec, Commissaria Smith said.
Destinations were the U.S., Covertible currency areas, Uruguay and open ports.
Cake sales were registered at 785 to 995 dlrs for March/April, 785 dlrs for May, 753 dlrs for Aug and 0.39 times New York Dec for Oct/Dec.
Buyers were the U.S., Argentina, Uruguay and convertible currency areas.
Liquor sales were limited with March/April selling at 2,325 and 2,380 dlrs, June/July at 2,375 dlrs and at 1.25 times New York July, Aug/Sept at 2,400 dlrs and at 1.25 times New York Sept and Oct/Dec at 1.25 times New York Dec, Commissaria Smith said.
Total Bahia sales are currently estimated at 6.13 mln bags against the 1986/87 crop and 1.06 mln bags against the 1987/88 crop.
Final figures for the period to February 28 are expected to be published by the Brazilian Cocoa Trade Commission after carnival which ends midday on February 27.
Reuter
</BODY></TEXT>
</REUTERS>
..
..

It may seem somewhat laborious to parse data in this manner especially when compared to the actual mechanics of the machine learning. However a large part of a quant researcher's day is often spent wrangling the data into a format usable by the analysis software! Hence it is worth getting some practice at it.

If we take a look at the topics file, **all-topics-strings.lc.txt**, we can see the following (we have removed most of it for brevity):

acq
alum
austdlr
austral
barley
bfr
bop
can

```
carcass
castor-meal
castor-oil
castorseed
citruspulp
cocoa
coconut
coconut-oil
coffee
copper
copra-cake
corn
...
...
silver
singdlr
skr
sorghum
soy-meal
soy-oil
soybean
stg
strategic-metal
sugar
sun-meal
sun-oil
sunseed
tapioca
tea
tin
trade
tung
tung-oil
veg-oil
wheat
wool
wpi
yen
zinc
```

We can see that there are 135 separate topics among the articles. This will make for quite a classification challenge!

At this stage we need to create what is known as a list of *predictor-response pairs*. This is a list of

two-tuples that contain the most appropriate class label and the raw document text, as two separate components. For instance, we wish to end up with a data structure, after parsing, which is similar to the following:

```
[
    ("cat", "It is best not to give them too much milk"),
    (
        "dog", "Last night we took him for a walk,
        but he had to remain on the leash"
    ),
    ..
    ..
    ("hamster", "Today we cleaned out the cage and prepared the sawdust"),
    ("cat", "Kittens require a lot of attention in the first few months")
]
```

To create this structure we will need to parse all of the Reuters files individually and add them to a grand corpus list. Since the file size of the corpus is rather low it will easily fit into available RAM on most modern laptops/desktops. However in production applications it is usually necessary to stream training data into a machine learning system and carry out partial fitting on each batch in an iterative manner. Our current goal is to now create the SGML Parser. To do this we will subclass Python's **HTMLParser** class to handle the specific tags in the Reuters dataset. Upon subclassing **HTMLParser** we override three methods: **handle_starttag**, **handle_endtag** and **handle_data**. They tell the parser what to do at the beginning of SGML tags, what to do at the closing of SGML tags and how to handle the data in between.

We also create two additional methods, **_reset** and **parse**, which are used to take care of internal state of the class and to parse the actual data in a chunked fashion, so as not to use up too much memory. Finally, a basic **__main__** function has been included to test the parser on the first set of data within the Reuters corpus:

```
from __future__ import print_function

import pprint
import re
try:
    from html.parser import HTMLParser
except ImportError:
    from HTMLParser import HTMLParser

class ReutersParser(HTMLParser):
    """
    ReutersParser subclasses HTMLParser and is used to open the SGML
    files associated with the Reuters-21578 categorised test collection.

    The parser is a generator and will yield a single document at a time.
```

Since the data will be chunked on parsing, it is necessary to keep some internal state of when tags have been "entered" and "exited". Hence the in_body, in_topics and in_topic_d boolean members.

```

"""
def __init__(self, encoding='latin-1'):
    """
    Initialise the superclass (HTMLParser) and reset the parser.
    Sets the encoding of the SGML files by default to latin-1.
    """
    HTMLParser.__init__(self)
    self._reset()
    self.encoding = encoding

def _reset(self):
    """
    This is called only on initialisation of the parser class
    and when a new topic-body tuple has been generated. It
    resets all off the state so that a new tuple can be subsequently
    generated.
    """
    self.in_body = False
    self.in_topics = False
    self.in_topic_d = False
    self.body = ""
    self.topics = []
    self.topic_d = ""

def parse(self, fd):
    """
    parse accepts a file descriptor and loads the data in chunks
    in order to minimise memory usage. It then yields new documents
    as they are parsed.
    """
    self.docs = []
    for chunk in fd:
        self.feed(chunk.decode(self.encoding))
        for doc in self.docs:
            yield doc
        self.docs = []
    self.close()

def handle_starttag(self, tag, attrs):
    """
    This method is used to determine what to do when the parser

```

```

comes across a particular tag of type "tag". In this instance
we simply set the internal state booleans to True if that particular
tag has been found.
"""
if tag == "reuters":
    pass
elif tag == "body":
    self.in_body = True
elif tag == "topics":
    self.in_topics = True
elif tag == "d":
    self.in_topic_d = True

def handle_endtag(self, tag):
    """
    This method is used to determine what to do when the parser
    finishes with a particular tag of type "tag".

    If the tag is a <REUTERS> tag, then we remove all
    white-space with a regular expression and then append the
    topic-body tuple.

    If the tag is a <BODY> or <TOPICS> tag then we simply set
    the internal state to False for these booleans, respectively.

    If the tag is a <D> tag (found within a <TOPICS> tag), then we
    append the particular topic to the "topics" list and
    finally reset it.
    """
    if tag == "reuters":
        self.body = re.sub(r'\s+', r' ', self.body)
        self.docs.append( (self.topics, self.body) )
        self._reset()
    elif tag == "body":
        self.in_body = False
    elif tag == "topics":
        self.in_topics = False
    elif tag == "d":
        self.in_topic_d = False
        self.topics.append(self.topic_d)
        self.topic_d = ""

def handle_data(self, data):
    """

```

```

        The data is simply appended to the appropriate member state
        for that particular tag, up until the end closing tag appears.
        """
        if self.in_body:
            self.body += data
        elif self.in_topic_d:
            self.topic_d += data

if __name__ == "__main__":
    # Open the first Reuters data set and create the parser
    filename = "data/reut2-000.sgm"
    parser = ReutersParser()

    # Parse the document and force all generated docs into
    # a list so that it can be printed out to the console
    doc = parser.parse(open(filename, 'rb'))
    pprint.pprint(list(doc))

```

At this stage we will see a significant amount of output that looks like this:

```

..
..
(['grain', 'rice', 'thailand'],
'Thailand exported 84,960 tonnes of rice in the week ended February 24, '
'up from 80,498 the previous week, the Commerce Ministry said. It said '
'government and private exporters shipped 27,510 and 57,450 tonnes '
'respectively. Private exporters concluded advance weekly sales for '
'79,448 tonnes against 79,014 the previous week. Thailand exported '
'689,038 tonnes of rice between the beginning of January and February 24, '
'up from 556,874 tonnes during the same period last year. It has '
'commitments to export another 658,999 tonnes this year. REUTER '),
(['soybean', 'red-bean', 'oilseed', 'japan'],
'The Tokyo Grain Exchange said it will raise the margin requirement on '
'the spot and nearby month for U.S. And Chinese soybeans and red beans, '
'effective March 2. Spot April U.S. Soybean contracts will increase to '
'90,000 yen per 15 tonne lot from 70,000 now. Other months will stay '
'unchanged at 70,000, except the new distant February requirement, which '
'will be set at 70,000 from March 2. Chinese spot March will be set at '
'110,000 yen per 15 tonne lot from 90,000. The exchange said it raised '
'spot March requirement to 130,000 yen on contracts outstanding at March '
'13. Chinese nearby April rises to 90,000 yen from 70,000. Other months '
'will remain unchanged at 70,000 yen except new distant August, which '
'will be set at 70,000 from March 2. The new margin for red bean spot '

```



```
'March rises to 150,000 yen per 2.4 tonne lot from 120,000 and to 190,000 '
'for outstanding contracts as of March 13. The nearby April requirement '
'for red beans will rise to 100,000 yen from 60,000, effective March 2. '
'The margin money for other red bean months will remain unchanged at '
'60,000 yen, except new distant August, for which the requirement will '
'also be set at 60,000 from March 2. REUTER '),
..
..
```

In particular, note that instead of having a single topic label associated with a document, we have multiple topics. In order to increase the effectiveness of the classifier, it is necessary to assign only a single class label to each document. However you will also note that some of the labels are actually geographic location tags, such as “japan” or “thailand”. Since we are concerned solely with topics and not countries we want to remove these before we select our topic. The particular method that we will use to carry this out is rather simple. We will strip out the country names and then select the first remaining topic on the list. If there are no associated topics we will eliminate the article from our corpus.

To remove the geographic tags and select the primary topic tag we can add the following code:

```
..
..

def obtain_topic_tags():
    """
    Open the topic list file and import all of the topic names
    taking care to strip the trailing "\n" from each word.
    """
    topics = open(
        "data/all-topics-strings.lc.txt", "r"
    ).readlines()
    topics = [t.strip() for t in topics]
    return topics

def filter_doc_list_through_topics(topics, docs):
    """
    Reads all of the documents and creates a new list of two-tuples
    that contain a single feature entry and the body text, instead of
    a list of topics. It removes all geographic features and only
    retains those documents which have at least one non-geographic
    topic.
    """
    ref_docs = []
    for d in docs:
```

```

    if d[0] == [] or d[0] == "":
        continue
    for t in d[0]:
        if t in topics:
            d_tup = (t, d[1])
            ref_docs.append(d_tup)
            break
    return ref_docs

if __name__ == "__main__":
    # Open the first Reuters data set and create the parser
    filename = "data/reut2-000.sgm"
    parser = ReutersParser()

    # Parse the document and force all generated docs into
    # a list so that it can be printed out to the console
    docs = list(parser.parse(open(filename, 'rb')))

    # Obtain the topic tags and filter docs through it
    topics = obtain_topic_tags()
    ref_docs = filter_doc_list_through_topics(topics, docs)
    pprint.pprint(ref_docs)

```

The output from this is as follows:

```

..
..
('acq',
'Security Pacific Corp said it completed its planned merger with Diablo '
'Bank following the approval of the comptroller of the currency. Security '
'Pacific announced its intention to merge with Diablo Bank, headquartered '
'in Danville, Calif., in September 1986 as part of its plan to expand its '
'retail network in Northern California. Diablo has a bank offices in '
'Danville, San Ramon and Alamo, Calif., Security Pacific also said. '
'Reuter '),
('earn',
'Shr six cts vs five cts Net 188,000 vs 130,000 Revs 12.2 mln vs 10.1 mln '
'Avg shrs 3,029,930 vs 2,764,544 12 mths Shr 81 cts vs 1.45 dlrs Net '
'2,463,000 vs 3,718,000 Revs 52.4 mln vs 47.5 mln Avg shrs 3,029,930 vs '
'2,566,680 NOTE: net for 1985 includes 500,000, or 20 cts per share, for '
'proceeds of a life insurance policy. includes tax benefit for prior qtr '
'of approximately 150,000 of which 140,000 relates to a lower effective '
'tax rate based on operating results for the year as a whole. Reuter '),

```

..
..

We are now in a position to pre-process the data for input into the classifier.

2. Vectorisation

At this stage we have a large collection of two-tuples, each containing a class label and raw body text from the articles. The obvious question to ask now is how do we convert the raw body text into a data representation that can be used by a (numerical) classifier?

The answer lies in a process known as **vectorisation**. Vectorisation allows widely-varying lengths of raw text to be converted into a numerical format that can be processed by the classifier. It achieves this by creating **tokens** from a string. A *token* is an individual word (or group of words) extracted from a document, using whitespace or punctuation as separators. This can include numbers from within the string as additional words. Once this list of tokens has been created they can be assigned an integer index identifier, which allows them to be listed.

Once the list of tokens has been generated the number of tokens within each document are counted. These tokens are then **normalised** to de-emphasise tokens that appear frequently within a document (such as "a", "the"). This process is known as the **Bag Of Words**. The Bag Of Words representation allows a vector to be associated with each document, each component of which is real-valued (i.e. $\in \mathbb{R}$) and represents the importance of tokens (i.e. words) appearing within that document.

Furthermore it means that once an entire corpus of documents has been iterated over (and thus all possible tokens have been assessed) the total number of separate tokens is known. Hence the length of the token vector for any document of any length in the training sample is also fixed and identical. This means that the classifier now has a set of features derived from the frequency of token occurrence. Each document token vector now represents a sample, or observation, x_j for the classifier. In essence the entire corpus can be represented as a large matrix. Each row of this matrix represents one of the documents and each column represents token occurrence within that document. This is the process of vectorisation.

Note: vectorisation does not take into account the relative positioning of the words within the document, just the frequency of instance. More sophisticated machine learning techniques do make use of this information to enhance the classification process.

3. Term-Frequency Inverse Document-Frequency

One of the major issues with vectorisation via the Bag Of Words representation is excessive noise in the form of **stop words**, such as "a", "the", "he", "she". These words provide little context to the document. Their relatively high frequency masks infrequently appearing words that do provide significant document context.

This motivates a transformation process known as **Term-Frequency Inverse Document-Frequency (TF-IDF)**. The TF-IDF value for a token increases proportionally to the frequency of the word in the document but is normalised by the frequency of the word in the corpus. This essentially reduces importance for words that appear a lot generally, as opposed to appearing a lot within a particular document. This is precisely what we need as words such as "a", "the" will have extremely high occurrence within the entire corpus, but the word "cat" may only appear often in a particular document. This would mean that we are giving "cat" a relatively higher strength than "a" or "the", for that document.

Hence we wish to combine the process of vectorisation with that of TF-IDF to produce a normalised matrix of document-token appearances. This will then be used to provide a list of features to the classifier upon which to train. The **TfidfVectorizer** class from Scikit-learn allows to realize such vectorization and transformation operations very easily. We can use this class to take our list of two-tuples representing class labels and raw document text, to produce both a vector of class labels and a sparse matrix, which represents the TF-IDF and vectorisation procedure applied to the raw text data.

Scikit-Learn classifiers take two separate data structures for training, namely y , the class label or response associated with a document, and, \mathbf{x} , the sparse TF-IDF vector associated with a document. We need to combine each of these length two tuples to create an ordered vector containing all y class labels and a similarly ordered matrix containing all TF-IDF vector rows, \mathbf{x} . The code to create these objects is given below:

```
..
from sklearn.feature_extraction.text import TfidfVectorizer
..
..

def create_tfidf_training_data(docs):
    """
    Creates a document corpus list (by stripping out the
    class labels), then applies the TF-IDF transform to this
    list.

    The function returns both the class label vector (y) and
    the corpus token/feature matrix (X).
    """
    # Create the training data class labels
    y = [d[0] for d in docs]

    # Create the document corpus list
    corpus = [d[1] for d in docs]

    # Create the TF-IDF vectoriser and transform the corpus
    vectorizer = TfidfVectorizer(min_df=1)
    X = vectorizer.fit_transform(corpus)
    return X, y

if __name__ == "__main__":
    # Open the first Reuters data set and create the parser
    filename = "data/reut2-000.sgm"
    parser = ReutersParser()

    # Parse the document and force all generated docs into
```

```

# a list so that it can be printed out to the console
docs = list(parser.parse(open(filename, 'rb'))))

# Obtain the topic tags and filter docs through it
topics = obtain_topic_tags()
ref_docs = filter_doc_list_through_topics(topics, docs)

# Vectorise and TF-IDF transform the corpus
X, y = create_tfidf_training_data(ref_docs)

```

At this stage we now have two components to our training data. The first, X , is a matrix of document-token occurrence. The second, Y , is a vector (which matches the ordering of the matrix) that contains the correct class labels for each of the documents. This is all we need to begin training and testing the support vector machine.

4. Training the Support Vector Machine

In order to train the support vector machine it is necessary to provide it with both a set of features (the X matrix) and a set of supervised training labels, in this case the Y classes vector. However we also need a means of evaluating the test performance of the classifier subsequent to its training phase. We discussed approaches for this in the prior chapter on Cross-Validation.

One question that arises here is what percentage to retain for training and what to use for testing. Clearly the more data retained for training the better the classifier will be because it will have seen more data. However more training data means less testing data and as such will lead to a poorer estimate of its true classification capability. In this section we will retain approximately 70-80% of the data for training and use the remainder for testing. A more sophisticated approach would be to use k-fold cross-validation.

Since the training-test split is such a common operation in machine learning, the developers of Scikit-Learn provided the `train_test_split` method to automatically create the split from a dataset provided (which we have already discussed in the previous chapter). Here is the code that provides the split:

```

..
..
from sklearn.cross_validation import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

```

The `test_size` keyword argument controls the size of the testing set, which in this case is 20%. The `random_state` keyword argument controls the random seed for selecting the random partitions into the training and test sets, which means the results should be identical on your implementation.

The next step is to actually create the support vector machine and train it. In this instance we are going to use the support vector classifier **SVC** class from Scikit-Learn. We give it the parameters `C=1000000.0`, `gamma="auto"` and choose a **radial kernel**.

The `C` parameter authorizes observations to violate the hyperplane separating the different classes (labels). In essence `C` is the parameter that governs the bias-variance trade-off for the SVC. A small value of `C` means a low bias, high variance situation. A large value of `C` means a high bias, low variance situation. This parameter is usually chosen via cross-validation. Here, `C=1000000.0` is a good value for such case but the reader can try other values to improve the current algorithm. The kernel is chosen as radial as it allows to create a separation such as a circle in 2 dimensions (similar to a cluster), and it can be generalized on higher dimensional hyperplanes. For more information on the SVC parameters, the reader is referred to the [RBF SVM Parameters](#) page of Scikit-Learn.

The following code imports the SVC class and then fits it on the training data:

```
from sklearn.svm import SVC

..
..

def train_svm(X, y):
    """
    Create and train the Support Vector Machine.
    """
    svm = SVC(C=1000000.0, gamma="auto", kernel='rbf')
    svm.fit(X, y)
    return svm

if __name__ == "__main__":
    # Open the first Reuters data set and create the parser
    filename = "data/reut2-000.sgm"
    parser = ReutersParser()

    # Parse the document and force all generated docs into
    # a list so that it can be printed out to the console
    docs = list(parser.parse(open(filename, 'rb'))))

    # Obtain the topic tags and filter docs through it
    topics = obtain_topic_tags()
    ref_docs = filter_doc_list_through_topics(topics, docs)

    # Vectorise and TF-IDF transform the corpus
    X, y = create_tfidf_training_data(ref_docs)

    # Create the training-test split of the data
    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.2, random_state=42
    )
```

```
# Create and train the Support Vector Machine
svm = train_svm(X_train, y_train)
```

Now that the SVM has been trained we need to assess its performance on the testing data.

5. Performance Metrics

The two main performance metrics that we will consider for this supervised classifier are the **hit-rate accuracy** and the **confusion matrix**. The former is simply the ratio of correct assignments to total assignments and is usually quoted as a percentage.

The confusion matrix goes into more detail and provides output on true-positives, true-negatives, false-positives and false-negatives. In a binary classification system, with a "true" or "false" class labelling, these characterise the rate at which the classifier correctly classifies an entity as true or false when it is, respectively, true or false, and also incorrectly classifies an entity as true or false when it is, respectively, false or true.

A confusion matrix need not be restricted to a binary classification situation. For multiple class groups (as in our situation with the Reuters dataset) we will have an $N \times N$ matrix, where N is the number of class labels (or document topics). Scikit-Learn has functions for calculating both the hit-rate and the confusion matrix of a supervised classifier. The former is a method on the classifier itself called **score**. The latter must be imported from the **metrics** library.

The first task is to create a predictions array from the **X_test** test-set. This will simply contain the predicted class labels from the SVM via the retained 20% test set. This prediction array is used to create the confusion matrix. Notice that the **confusion_matrix** function takes both the **pred** predictions array and the **y_test** correct class labels to produce the matrix. In addition we create the hit-rate by providing score with both the **X_test** and **y_test** subsets of the dataset:

```
..
..
from sklearn.metrics import confusion_matrix
..
..
if __name__ == "__main__":
    ..
    ..

    # Create and train the Support Vector Machine
    svm = train_svm(X_train, y_train)

    # Make an array of predictions on the test set
    pred = svm.predict(X_test)

    # Output the hit-rate and the confusion matrix for each model
    print(svm.score(X_test, y_test))
    print(confusion_matrix(pred, y_test))
```

The output of the code is as follows:

```
0.660194174757
[[21 0 0 0 2 3 0 0 0 1 0 0 0 0 1 1 1 0 0]
 [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [ 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [ 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [ 0 0 0 0 4 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [ 0 1 0 0 1 26 0 0 0 1 0 1 0 1 0 0 0 0]
 [ 0 0 0 0 0 0 2 0 0 0 0 0 0 0 0 0 0 1]
 [ 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0]
 [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [ 0 0 0 0 1 0 0 0 0 3 0 0 0 0 0 0 0 0]
 [ 3 0 0 1 2 2 3 0 1 1 6 0 1 0 0 0 2 3]
 [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [ 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0]
 [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0]
 [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]]
```

Thus we have a 66% classification accuracy, with a confusion matrix that has entries mainly on the diagonal (i.e. the correct assignment of class label). Notice that since we are only using a single file from the Reuters set (number 000), we are not going to see the entire set of class labels. Hence our confusion matrix is smaller in dimension than would be the case if we had used the full dataset.

In order to make use of the full dataset we can modify the `--main--` function to load all 21 Reuters files and train the SVM on the full dataset. We can output the full hit-rate performance. The confusion matrix output has not been included as it becomes large for the total number of class labels within all documents. Note that this will take some time, approximately 30-45 seconds, to run.

For the full corpus, the hit rate provided is 83.6%:

```
0.835971855761
```

There are plenty of ways to improve on this figure. In particular we can perform a **Grid Search Cross-Validation**, which is a means of determining the optimal parameters for the classifier that will achieve the best hit-rate (or other metric of choice). This is left for the reader as an exercise.

References

- [1] Wikipedia. Standard generalized markup language, 2015. URL https://en.wikipedia.org/wiki/Standard_Generalized_Markup_Language.

- [2] Python html parser, 2016. URL <https://docs.python.org/2/library/htmlparser.html>.